

Ch5: Functions and Recursion

Global Functions

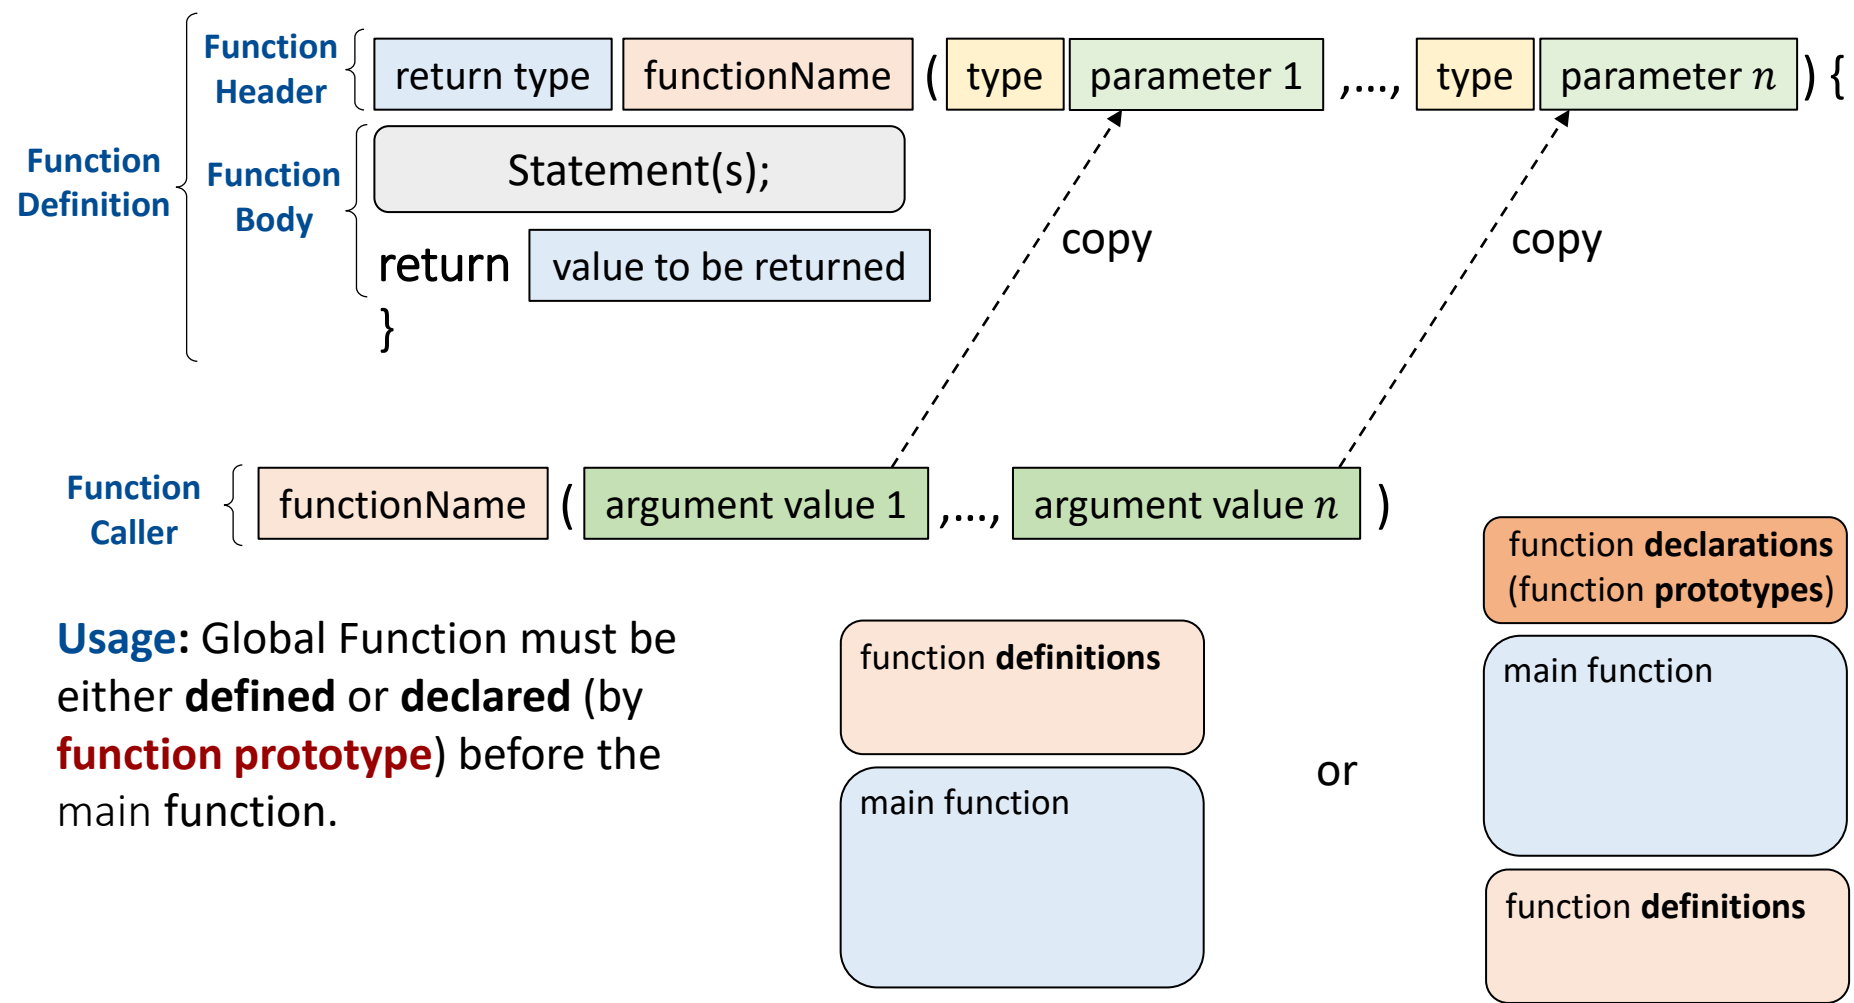
Global Functions

Functions allow you to modularize/divide a program by separating its tasks into self-contained units.

Advantages:

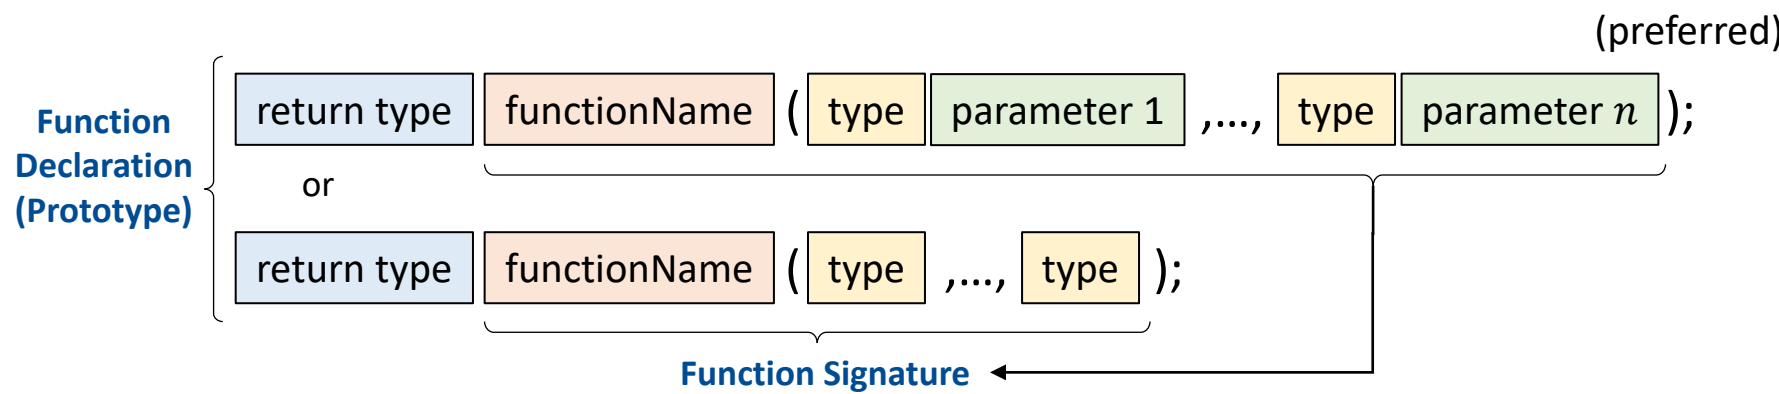
- Software reuse
 - Avoiding code repetition.
 - Making the program easier to test, debug, and maintain.
- ❖ The functions that are not members of a class are called **Global Functions** (e.g., main function or all functions in the [<cmath>](#) header that perform common mathematical calculations).

Function Definition and Function Caller



Function Declaration/Prototype and Function Signature


A **function prototype/declaration** is the same as the first line of the corresponding function definition but ends with a required semicolon.




Remarks

- If a function does not require any parameter to perform a task, its parameter list must be empty as ().
- The **return type** specifies the type of data the function returns to its caller after performing its task. If a function does not return any information to its caller, its type must be **void**.
- The **argument types** in the function call must be **consistent** (not necessarily identical) with the types of the corresponding parameters in the function's definition, otherwise, the compiler attempts to perform Implicit Type Conversion (Narrowing Conversions or Promotion) to convert the arguments to those types. Narrowing conversions can result in incorrect values.

- Commas in function calls are not comma operators. The order of evaluation of a function's arguments is not specified by the C++ standard. Thus, different compilers can evaluate function arguments in different orders.



```
#include <iostream>
double addDouble(double x, double y) {
    return x + y;
}
int addInt(int x, int y) {
    return x + y;
}
int main() {
    std::cout << "3.4 + 4 = " << addDouble(3.4, 4) << "\n"; // promotion in 4
    std::cout << "3.4 + 4 = " << addInt(3.4, 4) << "\n"; // narrowing in 3.4
}
```



Sample Program: maximum Function

A user-defined function called maximum that returns the largest of its three int arguments.

```
// maximum function with a function prototype.
#include <iostream>

int maximum(int x, int y, int z); // function prototype

int main() {
    std::cout << "Enter three integer values: ";
    int int1, int2, int3;
    std::cin >> int1 >> int2 >> int3;

    // invoke maximum
    std::cout << "The maximum integer value is: "
        << maximum(int1, int2, int3) << std::endl;
}
```

```
// returns the largest of three integers
int maximum(int x, int y, int z) {
    int maximumValue{x}; // assume x is the largest to start

    // determine whether y is greater than maximumValue
    if (y > maximumValue) {
        maximumValue = y;
    } // make y the new maximumValue

    // determine whether z is greater than maximumValue
    if (z > maximumValue) {
        maximumValue = z; // make z the new maximumValue
    }

    return maximumValue;
}
```



Function Declarations and Definitions in Separate Files (Best Practices)

- You should never put function definitions in header files (except **template functions**). Header files should contain function declarations/prototypes.
- Function definitions can be in a separate .cpp file. Each .cpp file should have #include directives for any functions that it calls.
- All these files (separate header containing function declarations and .cpp containing function definitions) must be into the project folder.

main.cpp

```
#include <iostream>
#include "Functions.h"

int main() {
    std::cout << "3 + 4 = " << add(3, 4) << '\n';
    std::cout << "3^2 + 4^2 = " << sumOfSquares(3, 4) << '\n';
}
```

Functions.h

```
int add(int x, int y);
int sumOfSquares(int x, int y);
int pow2(int x);
```

MyFunctions.cpp

```
#include "Functions.h"
int add(int x, int y) {
    return x + y;
}
int sumOfSquares(int x, int y) {
    return pow2(x) + pow2(y);
}
int pow2(int x) {
    return x*x;
}
```


const Function Parameters

- Function parameters can be made constants via the **const** keyword, without initialization.
- Making a function parameter constant inform the compiler that the parameter's value is not accidentally changed inside the function.
- The const qualifier should be used to enforce the principle of least privilege. Using this principle to properly design software can greatly reduce debugging time and improper side effects and can make a program easier to modify and maintain.


```
#include <iostream>
int add(const int x, int y) {
    // x++; // error
    ++y;
    return x + y;
}
int main() {
    std::cout << "Result: " << add(3, 4) << "\n";
}
```

Scope Rules

Block Scope and Local Variable

- The portion of a program where an identifier can be used is known as its **Scope**.
- Identifiers declared inside a block ({}) have **block scope**, which begins at the identifier's declaration and ends at the terminating right brace (}) of the enclosing block (the blocks can appear in **all control structures and functions**).
- **Local variables** have block scope, and can be referenced only in that block and in blocks nested within that block (inner blocks).

```
#include <iostream>
int main() {
    int y{5}; // local variable to main
    std::cout << "y = " << y << "\n";
    {
        int z{7}; // local variable in nested block
        // "y" is accessible here
        std::cout << "y = " << y << "\n" << "z = " << z << "\n";
    }
    // "z" is not accessible here
}
```



- Variables declared in a particular function's body are **local variables** which can be used only in that function. When a function terminates, the values of its local variables are lost.
- Parameters of a function also are **local variables** of that function.

static Local Variable


Local variables in a function can be declared **static**. Such variables not only have block scope, but also retains its value when the function returns to its caller. The next time the function is called, the static local variable contains the value it had when the function last completed execution.

Note: All static local variables of numeric types are initialized to zero by default.

```
#include <iostream>

void demo() {
    static int count; // static variable (it is initialized to zero by default)
    ++count; // value is updated and will be carried to next function calls
    std::cout << count << "\n";
}


int main() {
    for (int i{0}; i < 5; ++i){
        demo();
    }
}
```



Global Namespace Scope: Global Variables

- An identifier declared outside any function (including main), or class definition has **global namespace scope**. Such an identifier is “known” in all functions from the point at which it is declared until the end of the file.
- Function definitions, function prototypes placed outside a function, class definitions, and global variables all have global namespace scope.

```
#include <iostream>
int y{5}; // global variable
int Foo(int x) {
    y++; // function can read & modify global variables!
    return x + y;
}
int main() {
    std::cout << "y = " << y << "\n";
    std::cout << "Function = " << Foo(2) << "\n";
    std::cout << "y = " << y << "\n";
}
```



Unary Scope Resolution Operator

A global variable can be accessed directly; however, C++ provides the **Unary Scope Resolution Operator** (::) to access a global variable when a local variable of the same name is in scope.

```
#include <iostream>
int y{5}; // global variable
int Foo(int x) {
    ::y++; // function can read & modify global variables!
    return x + ::y;
}
int main() {
    int y{0}; // local variable
    std::cout << "y = " << y << "\n";
    std::cout << "y = " << ::y << "\n";
    std::cout << "Function = " << Foo(2) << "\n";
    std::cout << "y = " << ::y << "\n";
}
```

Best Practices:

- In general, variables should be declared in the narrowest scope in which they need to be accessed. Thus, global variables should be avoided.
- Avoid using variables of the same name for different purposes in a program. Although this is allowed in various circumstances, it can lead to errors.
- Always use (::) to refer to global variables (even if there is no collision with a local-variable name) to make it clear that you are intending to access a global variable rather than a local variable.

Recursion

Recursion

A recursive function is a function that calls itself, either directly, or indirectly (through another function).

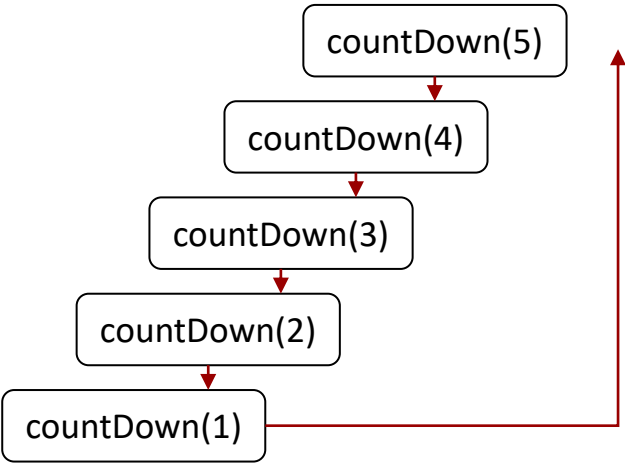
Recursive function calls generally work just like normal function calls. However, they need a recursive termination condition to stop calling itself (usually using an if statement), or they will run forever (actually, until the call stack runs out of memory).

```
#include <iostream>

void countDown(int i) {
    std::cout << "push " << i << '\n';
    if (i > 1) { // termination condition
        countDown(i - 1);
    }
    std::cout << "pop " << i << '\n';
}

int main() {
    countDown(5);
}
```

push 5
push 4
push 3
push 2
push 1
pop 1
pop 2
pop 3
pop 4
pop 5



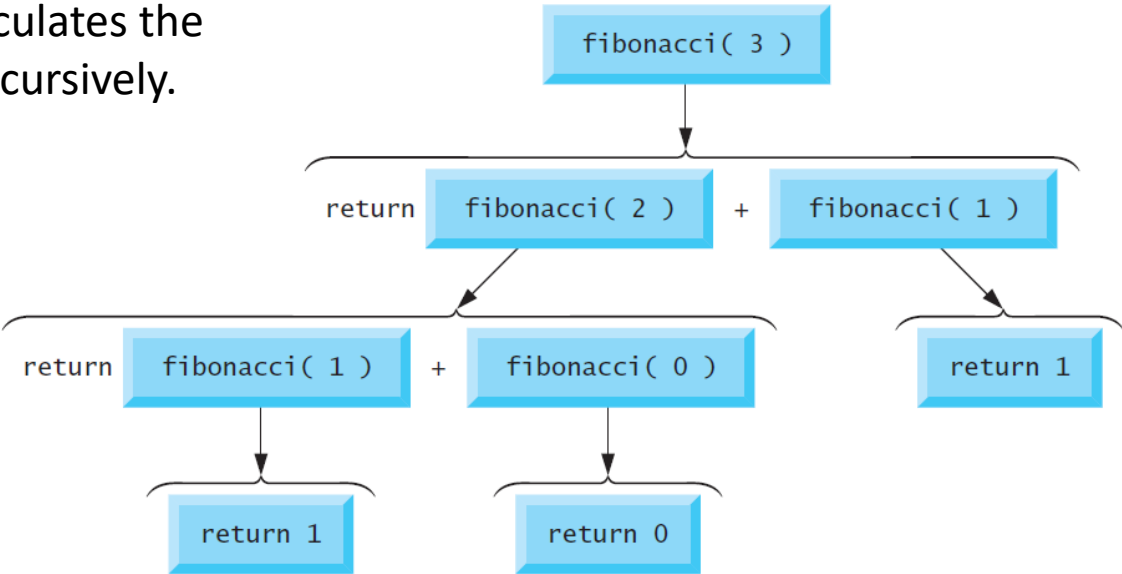
Example: Fibonacci Numbers

Fibonacci numbers are defined mathematically as:

$$f(n) = \begin{cases} 0 & \text{if } n = 0 \\ 1 & \text{if } n = 1 \\ f(n - 1) + f(n - 2) & \text{if } n > 1 \end{cases}$$

0, 1, 1, 2, 3, 5, 8, 13, 21, ...
 (begins with 0 and 1 and each subsequent Fibonacci number is the sum of the previous two Fibonacci numbers)

Write a function that calculates the *n*th Fibonacci number recursively.




Example: Fibonacci Numbers

```
#include <iostream>
unsigned long fibonacci(unsigned long number); // function prototype
int main() {
    // calculate the fibonacci values of 0 through 10
    for (unsigned int n{0}; n <= 10; ++n) {
        std::cout << "fibonacci(" << n << ") = " << fibonacci(n) << std::endl;
    }

    // display higher fibonacci values
    std::cout << "fibonacci(30) = " << fibonacci(30) << std::endl;
}

// recursive function fibonacci
unsigned long fibonacci(unsigned long number) {
    if ((0 == number) || (1 == number)) { // base case (termination condition)
        return number;
    }
    else { // recursion step
        return fibonacci(number - 1) + fibonacci(number - 2);
    }
}
```



Recursion vs Iteration

- Any problem that can be solved recursively can also be solved iteratively (non-recursively).
- A recursive approach is normally chosen when the recursive approach results in a program that is easier to understand and debug. However, recursive calls take time and consume additional memory.

```
#include <iostream>
unsigned long fibonacci (unsigned long number); // function prototype
int main() {
    // calculate the fibonacci values of 0 through 10
    for (unsigned int n{0}; n <= 10; ++n) {
        std::cout << "fibonacci(" << n << ") = " << fibonacci(n) << std::endl;
    }
    // display higher fibonacci values
    std::cout << "fibonacci(30) = " << fibonacci(30) << std::endl;
}

unsigned long fibonacci (unsigned long number) {
    int t1{0}, t2{1}, nextTerm{0};
    // The first two terms.
    if(number == 0)
        return t1;
    if(number == 1)
        return t2;
    for (int n{2}; n <= number; ++n) {
        nextTerm = t1 + t2;
        t1 = t2;
        t2 = nextTerm;
    }
    return nextTerm;
}
```



Function Overloading

Function Overloading

Function Overloading is used to create several functions of the **same name** (that perform similar tasks), as long as they have **different signatures**.

- The C++ compiler selects the proper function to call by examining the number, types, and order of the arguments in the call.

```
#include <iostream>

int add(int x, int y);
double add(double x, double y);
int add(int x, int y, int z);

int main() {
    std::cout << add(1, 2) << "\n"; // calls int version
    std::cout << add(1.2, 2.3) << "\n"; // calls double version
    // std::cout << add(1, 2.3) << "\n"; // error: call of overloaded
    // 'add(int, double)' is ambiguous
    std::cout << add(1, 2, 3) << "\n"; // calls int version
}
```

```
// function add for int values
int add(int x, int y) {
    return x + y;
}

// function add for double values
double add(double x, double y) {
    return x + y;
}

// function add for int values
int add(int x, int y, int z) {
    return x + y + z;
}
```

Default Arguments

Default Arguments

A function can have a parameter with default argument. When making a function call, the caller can optionally provide an argument for any function parameter that has a **Default Argument**.

- If the caller provides an argument, the value of the argument in the function call is used.
- If the caller does not provide an argument, the value of the default argument is used.

```
#include <iostream>

int add(int x, int y=4){ // 4 is the default argument
    return x + y;
}

int main() {
    std::cout << "x + y = " << add(1, 2) << '\n'; // y will use user-supplied argument 2
    std::cout << "x + y = " << add(3) << '\n'; // y will use default argument 4
}
```

Note: The equals sign must be used to specify a default argument. Using brace initialization will not work:

```
int add(int x, int y{4}) // compile error
```


Multiple Default Arguments

A function can have multiple parameters with default arguments:

```
#include <iostream>

void print(int x = 10, int y = 20, int z = 30) {
    std::cout << "Values: " << x << " " << y << " " << z << "\n";
}

int main() {
    print(1, 2, 3); // all explicit arguments
    print(1, 2); // rightmost argument defaulted
    print(1); // two rightmost arguments defaulted
    print(); // all arguments defaulted
}
```



Note: Default arguments can only be provided for the rightmost parameters.

```
void print(int x = 10, int y = 20, int z) // not allowed
```

Because C++ (as of C++20) does not support a function call syntax such as `print(.,3)`

Default Arguments and Function Declaration

Once declared, a default argument can not be redeclared. Therefore, the default argument can be declared in either the function declaration or the function definition, but not both.

```
#include <iostream>

int add(int x, int y=4); // function declaration (function prototype)

int main() {
    std::cout << "x + y = " << add(1, 2) << "\n"; // y will use user-supplied argument 2
    std::cout << "x + y = " << add(3) << "\n"; // y will use default argument 4
}

int add(int x, int y) { // function definition
    return x + y;
}
```

Best practice: Declare the default argument in the function declaration and not in the function definition (particularly if it's in a header file).

Default Arguments and Function Overloading

Functions with default arguments may be overloaded. However, such functions can lead to potentially **ambiguous function calls**.

```
#include <iostream>

int add(int x){
    return x;
}

int add(int x, int y = 4){
    return x + y;
}

double add(int x, double y = 4.6){
    return x + y;
}

int main() {
    std::cout << "add(1, 2) = " << add(1, 2) << "\n"; // will resolve to add(int, int)
    std::cout << "add(1, 2.5) = " << add(1, 2.5) << "\n"; // will resolve to add(int, double)
    std::cout << "add(1) = " << add(1) << "\n"; // ambiguous function call
}
```

```
#include <iostream>

void print(std::string string) {
    std::cout << string << "\n";
}

void print(char ch = ' ') {
    std::cout << ch << "\n";
}

int main() {
    print("Hello, World"); // resolves to print(std::string)
    print('a'); // resolves to print(char)
    print(); // resolves to print(char)
}
```

Function Templates

Introduction to Function Templates

Consider the following code. To calculate the addition of two integer and floating-point numbers, we must create overloaded versions of function `add` with parameters of type `int` and `double` (similarly, `long`, `long double`, and even new types that you've created).

```
#include <iostream>
int add(int x, int y) {
    return x + y;
}

double add(double x, double y) {
    return x + y;
}

int main() {
    std::cout << add(1, 2) << "\n";
    std::cout << add(1.2, 2.2) << "\n";
}
```

Overloaded functions are normally used to perform similar operations that involve **different** program logic on different data types. If the program logic and operations are **identical** for each data type, we can instead create and maintain a single version of function (**template**) that can work with arguments of any type.

Function Template

```
#include <iostream>

template <typename T> // this is the template parameter declaration
T add(T x, T y) { // this is the function template definition
    return x + y;
}

int main() {
    std::cout << add<int>(1, 2) << "\n"; // instantiates and calls function add<int>(int, int)
    std::cout << add<double>(1, 2.2) << "\n"; // instantiates and calls function add<double>(double, double)
                                           // 1 will be implicitly converted to double
}
```

- When we create our function template, we use **placeholder types** (also called **type template parameters**) for fundamental types or user-defined types. These placeholders are used to specify the parameter types, return types, or types used in the function body that we want to be specified later.
- The type in angled brackets (called a template argument) specifies the actual type that will be used in place of template type (T) by the compiler.
- **Best Practice:** Use a single capital letter (starting with T) to name your type template parameters (e.g., T, U, V, ...).

Template Argument Deduction

In cases where the type of the arguments match the actual type we want, we do not need to specify the actual type; instead, we can use template argument deduction to have the compiler deduce the actual type that should be used from the argument types in the function call.

```
#include <iostream>

template <typename T> // this is the template parameter declaration
T add(T x, T y) { // this is the function template definition
    return x + y;
}

int main() {
    std::cout << add<>>(1, 2) << "\n"; // calls function add<int>(int, int)
    std::cout << add(1, 2) << "\n"; // calls function add<int>(int, int)
    std::cout << add(1.0, 2.2) << "\n"; // calls function add<double>(double, double)
    // std::cout << add(1, 2.2) << "\n"; // syntax error: no matching function for call to "add(int, double)"
}
```

Best Practice: Favor the normal function call syntax when using function templates.

Implicit Type Conversion

```
#include <iostream>

template <typename T> // this is the template parameter declaration
T add(T x, T y) { // this is the function template definition
    return x + y;
}

double add2(double x, double y) {
    return x + y;
}

int main() {
    std::cout << add<double>(1, 2.2) << "\n"; // calls function add<double>(double, double)
                                                // 1 will be implicitly converted to double
    std::cout << add(1.0, 2.2) << "\n"; // calls function add<double>(double, double)
    // std::cout << add(1, 2.2) << "\n"; // syntax error: no matching function for call to "add(int, double)"
    std::cout << add(static_cast<double>(1), 2.2) << "\n"; // calls function add<double>(double, double)
    std::cout << add2(1, 2.2) << "\n"; // int argument will be implicitly converted to a double
}
```

Note: Implicit type conversions can be done when calling non-template functions. However, when performing template argument deduction, implicit type conversion are not allowed.

Functions Template Declaration

Functions templates can be declared (by function prototypes) before function main and defined after function main, but this definition cannot be placed in a separate source (.cpp) file. The best practice is to put all your function templates in a header (.h) file (without additional declarations) instead of a separate source (.cpp) file, and then, #included wherever needed.

```
#include <iostream>

template <typename T>
T addOne(T x); // function template declaration

int main() {
    std::cout << addOne(1) << "\n";
    std::cout << addOne(2.3) << "\n";
}

template <typename T>
T addOne(T x) { // function template definition
    return x + 1;
}
```

≡

addone.h

```
template <typename T>
T addOne(T x) { // function template definition
    return x + 1;
}
```

```
#include <iostream>
#include "addone.h"

int main() {
    std::cout << addOne(1) << "\n";
    std::cout << addOne(2.3) << "\n";
}
```



Example: maximum Function Template

```
#include <iostream>
#include "maximum.h"

int main() {
    std::cout << "Input three integer values: ";
    int int1, int2, int3;
    std::cin >> int1 >> int2 >> int3;
    std::cout << "The maximum integer value is: "
        << maximum(int1, int2, int3);

    std::cout << "\n\nInput three double values: ";
    double double1, double2, double3;
    std::cin >> double1 >> double2 >> double3;
    std::cout << "The maximum double value is: "
        << maximum(double1, double2, double3);

    std::cout << "\n\nInput three characters: ";
    char char1, char2, char3;
    std::cin >> char1 >> char2 >> char3;
    std::cout << "The maximum character value is: "
        << maximum(char1, char2, char3) << std::endl;
}
```

maximum.h

```
template <typename T>
T maximum(T value1, T value2, T value3) {
    T maximumValue{value1};

    if (value2 > maximumValue) {
        maximumValue = value2;
    }

    if (value3 > maximumValue) {
        maximumValue = value3;
    }

    return maximumValue;
}
```




Function Templates with Non-Template Parameters

It is possible to create function templates that have both template parameters and non-template parameters. The type template parameters can be matched to any type, and the non-template parameters work like the parameters of normal functions.

```
#include <iostream>

// T is a type template parameter
// y is a non-template parameter
template <typename T>
double add(T x, double y) {
    return x + y;
}

int main() {
    std::cout << add(1, 3.4) << "\n"; // matches add(int, double)
    std::cout << add(1.2, 3.4) << "\n"; // matches add(double, double)
}
```



Function Templates with Multiple Template Type Parameters

Rather than using one template type parameter T , we can use more (e.g., T, U, V, \dots) to resolve the types independently.

```
#include <iostream>
template <typename T, typename U> // using two template type parameters T and U
T add(T x, U y) { // x can resolve to type T, and y can resolve to type U
    return x + y;
}
int main() {
    std::cout << add(2.2, 1) << "\n"; // works fine, prints 3.2
    std::cout << add(1, 2.2) << "\n"; // narrowing conversion problem! prints 3
}
```

In the cases that narrowing conversion is an issue, we can use **auto** keyword for return type to let the compiler deduce what the return type should be from the return statement.

```
#include <iostream>
template <typename T, typename U>
auto add(T x, U y) {
    return x + y;
}
int main() {
    std::cout << add(2.2, 1) << "\n"; // prints 3.2
    std::cout << add(1, 2.2) << "\n"; // prints 3.2
}
```

Abbreviated Function Templates

C++20 introduces a new use of the **auto** keyword: When the auto keyword is used as a parameter type in a normal function, the compiler will automatically convert the function into a function template with each auto parameter becoming an independent template type parameter. This method for creating a **function template** is called an **abbreviated function template** that makes your code more concise and readable.

```
#include <iostream>

template <typename T, typename U>
auto add(T x, U y) {
    return x + y;
}

int main() {
    std::cout << add(2.2, 1) << "\n"; // prints 3.2
    std::cout << add(1, 2.2) << "\n"; // prints 3.2
}
```

≡

```
#include <iostream>

auto add(auto x, auto y) {
    return x + y;
}

int main() {
    std::cout << add(2.2, 1) << "\n"; // prints 3.2
    std::cout << add(1, 2.2) << "\n"; // prints 3.2
}
```

Lvalue Reference

Lvalues and Rvalues

Lvalues (left values) expressions are those that evaluate to variables or other identifiable objects that persist beyond the end of execution of the expression. They can be used on an assignment operator's left or side and come in two subtypes:

- **modifiable lvalue** whose value can be modified,
- **non-modifiable lvalue** whose value cannot be modified.

Rvalues (right values) expressions are those that evaluate to literals or the returned value of functions and operators that are discarded at the end of execution of the expression. They can be used on only an assignment operator's right side, but not vice versa.

Note: An assignment operation requires the left operand of the assignment to be a modifiable lvalue expression, and the right operand to be an rvalue expression, thus, $x = 5$ is valid but $5 = x$ is not.

Lvalues and Rvalues


```
#include <iostream>

int return5() {
    return 5;
}

int main() {
    int x{ 5 }; // 5 is an rvalue expression
    const double d{ 1.2 }; // 1.2 is an rvalue expression

    int y{ x }; // x is a modifiable lvalue expression
    const double e { d }; // d is a non-modifiable lvalue expression
    int z{ return5() }; // return5() is an rvalue expression (since the result is returned by value)

    int w{ x + 1 }; // x + 1 is an rvalue expression
    int q{ static_cast<int>(d) }; // the result of static casting d to an int is an rvalue expression
}
```



Lvalue Reference (or Reference)

A **Lvalue Reference** (commonly called a **Reference**) is an alias for an existing object (or variable). Once a reference has been defined, any operation on the reference is applied to the object (or variable) being referenced.

- To declare an lvalue reference type, an ampersand (&) is used in the type declaration.

```
int    // a normal int type
int&   // an lvalue reference to an int object
double& // an lvalue reference to a double object
```

- To create an lvalue reference variable, we simply define a variable with an lvalue reference type to read and modify the value of the object being referenced.

```
#include <iostream>
int main() {
    int x{5}; // x is a normal integer variable
    int& ref{x}; // ref is an lvalue reference variable that can now be used as an alias for variable x
    std::cout << x << ", " << ref << "\n"; // prints 5, 5 the value of x and the value of x via ref
    x = 6; // x now has value 6
    std::cout << x << ", " << ref << "\n"; // prints 6, 6 the value of x and the value of x via ref
    ref = 7; // the object being referenced (x) now has value 7
    std::cout << x << ", " << ref << "\n"; // prints 7, 7 the value of x and the value of x via ref
}
```


* The value of x can be changed through either x or ref.



Lvalue Reference (or Reference)

- Lvalue references must be **bound** to (or initialized with) a **modifiable lvalue**.
- Lvalue references cannot be bound to **non-modifiable lvalues** or **rvalues** (because it would allow us to modify a const variable through the non-const reference).
- The **type** of the reference must match the type of the referent (modifiable lvalue).
- Once initialized, a reference cannot be reassigned as aliases to other variables.
- Reference variables follow the same scoping and duration rules that normal variables do.

```
int main() {  
  
    int x{ 5 };  
    int& ref{ x }; // valid: lvalue reference bound to a modifiable lvalue  
  
    const int y{ 5 };  
    int& ref1{ y }; // invalid: can't bind to a non-modifiable lvalue  
    int& ref2{ 0 }; // invalid: can't bind to an r-value  
  
    double z{ 6.0 };  
    int& ref3{ z }; // invalid: reference to int cannot bind to double variable  
}
```



Lvalue Reference to const

- By using the `const` keyword when declaring an lvalue reference, we can create a reference to **non-modifiable lvalues** to access but not modify them.
- Lvalue reference to `const` can also bind to a **modifiable lvalue**. In such a case, we can use the reference to access the lvalue, but because reference is `const`, we can not modify the value of the lvalue through the reference. However, we still can modify the value of the lvalue directly.

```
#include <iostream>
int main() {
    const int x { 5 }; // x is a non-modifiable lvalue
    const int& ref1 { x }; // okay: ref1 is a an lvalue reference to a const value

    std::cout << ref1 << "\n"; // okay: we can access the const object
    ref1 = 6; // error: we can not modify a const object

    int y { 5 }; // y is a modifiable lvalue
    const int& ref2 { y }; // okay: we can bind a const reference to a modifiable lvalue

    std::cout << ref2 << "\n"; // okay: we can access the object through our const reference
    ref2 = 7; // error: we can not modify an object through a const reference
    y = 6; // okay: y is a modifiable lvalue, we can still modify it through the original identifier
}
```

Best Practice: Favor lvalue references to `const` unless you need to modify the object being referenced through the reference.

Lvalue Reference to const

We can initialize an lvalue reference to const with an rvalue. When this happens, a temporary object is created and initialized with the rvalue, and the reference to const is bound to that temporary object. The lifetime of the temporary object is extended to match the lifetime of the reference.

```
#include <iostream>
int main() {
    const int& ref{ 5 }; // A temporary object holding value 5 is created
                        // and ref is bound to it (5 is an rvalue)

    std::cout << ref << "\n"; // Prints 5
} // Both ref and the temporary object die here
```

Summary:

- Lvalue references can only bind to modifiable lvalues.
- Lvalue references to const can bind to modifiable lvalues, non-modifiable lvalues, and rvalues. This makes them a much more flexible type of reference.

Passing Arguments by Reference

Passing Arguments to Functions

Two ways to pass arguments to functions are **pass-by-value** and **pass-by-reference**.

- With **pass-by-value**:
 - A **copy** of the argument's value is made and passed to the called function.
 - Changes to the copy in the called function do not affect the original variable's value in the caller.
 - **Disadvantage** is that if a large data item is being passed, copying that data can take a considerable amount of execution time and memory space.
- With **pass-by-reference**:
 - The caller gives the called function the ability to **access the caller's data directly**.
 - It is **good for performance** reasons, because it can eliminate the pass-by-value overhead of copying large amounts of data.

Pass-by-Value

```
#include <iostream>
#include <string>
void printValue(int y) {
    std::cout << y << '\n';
} // y is destroyed here
void printValue(std::string y) {
    std::cout << y << '\n';
} // y is destroyed here
int main() {
    int x{2};
    printValue(x); // x is passed by value (copied) into parameter y (inexpensive)
    std::string s{"Hello, world!"}; // s is a std::string
    printValue(s); // s is passed by value (copied) into parameter y (expensive)
}
```

- **Fundamental types** are **cheap** to copy, however, most of the types provided by the standard library (such as `std::string`) are **class types** that are usually **expensive** to copy.
- One way to avoid making an expensive copy of an argument when calling a function and make the program more efficient is to use **pass-by-reference** instead of pass-by-value.

Pass-by-Reference

(reference to non-const)

When using pass by reference, we declare a function parameter as a reference type. When the function is called, each reference parameter is bound to the appropriate argument in the caller. Because the reference acts as an alias for the argument, no copy of the argument is made.

When using pass by reference to non-const,

- any changes made to the reference parameter in the function will affect the original value of the argument.
- only modifiable lvalue arguments are acceptable.

```
#include <iostream>
#include <string>
void printValue(std::string& y) { // y is bound to a modifiable lvalue (s)
    std::cout << y << '\n';
} // y is destroyed here
void printValue(int& y) { // y is bound to a modifiable lvalue
    std::cout << y++ << '\n'; // this modifies the actual object x
} // y is destroyed here
int main() {
    std::string s{"Hello, world!"}; // s is a modifiable lvalue
    printValue(s); // s is passed by reference
    int x{ 5 }; // x is a modifiable lvalue
    printValue(x); // x is passed by reference
    std::cout << x << '\n'; // x has been modified
    const int z{ 5 }; // z is a non-modifiable lvalue
    printValue(z); // error: z is a non-modifiable lvalue
    printValue(5); // error: 5 is an rvalue
}
```



Pass-by-Reference

(reference to const)

A reference to const can (1) bind to modifiable lvalues, non-modifiable lvalues, and rvalues (i.e., any type of argument) (2) guarantee that the function can not change the value being referenced (in most cases, we don't want our functions modifying the value of arguments).


```
#include <iostream>

void printValue(const int& y) { // y is a const reference
    std::cout << y << '\n';
    // ++y; // not allowed: y is const
}

int main() {
    int x{ 5 };
    printValue(x); // ok: x is a modifiable lvalue

    const int z{ 5 };
    printValue(z); // ok: z is a non-modifiable lvalue

    printValue(5); // ok: 5 is a literal rvalue
}
```



Mixing Pass-by-Value and Pass-by-Reference

A function with multiple parameters can determine whether each parameter is passed by value or passed by reference individually.

As always, the function prototype and header must agree.

```
#include <iostream>
#include <string>
void printValue(int a, int& b, const std::string& c);
int main() {
    int x{ 5 };
    const std::string s{ "Hello, world!" };
    printValue(15, x, s);
}
void printValue(int a, int& b, const std::string& c) {
    std::cout << a << ", " << b << ", " << c << "\n";
}
```

Best Practices:

- Favor passing by const reference over passing by non-const reference unless you have a specific reason to do otherwise (e.g., the function needs to change the value of an argument).
- Prefer pass-by-value for objects that are cheap to copy (e.g., fundamental types, enumerated types), and pass-by-const-reference for objects that are expensive to copy (e.g., class types including `std::string`, `std::array`, `std::vector`). If you're not sure whether an object is cheap or expensive to copy, favor pass-by-const-reference.